AIPI AI Survey
Wave 2

YouGov

Sample	Online sample of 1,118 voters fielded from September 02 to September 06, 2023.
Margin of Error	±3.2%

1. How important do you think the issue of Artificial Intelligence (AI) is in public policy?

Extremely important	43%
Somewhat important	35%
Not very important	7%
Not at all important	5%
Don't know	9%
Totals	99%
N	**1,118**

2. Do you [agree or disagree] that the following should be important goals of AI (Artificial Intelligence) policy?

	Strongly agree	Somewhat agree	Somewhat disagree	Strongly disagree	Not sure
Reducing the spread of dangerous AI	64%	22%	4%	2%	9%
Limiting dangers during model creation	53%	26%	5%	3%	12%
Reducing the dangers of very powerful AI which may be built in the next few years	57%	25%	5%	2%	11%
Restricting the speed of AI capability increase	38%	27%	11%	4%	20%
Reducing the dangers of AI today	58%	24%	5%	2%	10%
Reducing misinformation	58%	21%	5%	5%	11%
Treating AI as an incredibly powerful and dangerous technology	52%	26%	8%	3%	11%
Preventing AI from quickly reaching superhuman capabilities	56%	20%	8%	4%	12%
Preventing AI from discriminating on gender or race	58%	23%	5%	5%	9%
Preventing political bias in AI	61%	22%	6%	3%	9%

3. Which goal of AI policy is more important?

Keeping dangerous models out of the hands of bad actors	65%
Providing the benefits of AI to everyone	22%
Don't know	13%
Totals	100%
N	**1,118**

4. There is a debate around limiting AI models we don't understand. Some policymakers say that we don't understand how AI operates and how it will respond to different situations. They claim this is dangerous as the unknown capabilities of models grow, and that we should restrict models we don't understand. Other policymakers say that we understand broadly how AI models operate and that they're just statistical models. They say that limiting models until we have a full understanding is unrealistic and will put us behind competitors like China. What do you think? Should we place limits on AI models we don't fully understand?

Yes, we should limit these models	63%
No, we should not limit these models	14%
Don't know	23%
Totals	100%
N	**1,118**

5. There is a debate around regulating AI model creation. Proponents of regulating model creation say that models can be dangerous when they are trained, even before they are released to the public. They point out that bad actors like Chinese hackers could steal the technology and that dangerous accidents could occur when testing out a powerful model. Opponents of regulating model creation say that the real danger occurs when the models are released to the public and used. They say that bad actors stealing models before they're released is rare, that AI causing accidents in a lab setting is unrealistic and regulating creation is impractical. What do you think? Should we regulate model creation?

Yes, we should regulate model creation	65%
No, we should not regulate model creation	12%
Don't know	23%
Totals	100%
N	**1,118**

6. There is a debate around the scope of AI policies and whether they should prepare for more powerful AI models in the future. Supporters of preparing for powerful models say that because models are advancing rapidly we should take a proactive approach, developing regulation that will limit models further as they become more powerful. They claim that if we don't limit how powerful models can get now, we'll move too slow to stop dangers. Opponents of preparing for powerful models say that we can't prepare for models that don't exist. They say that all we can do is regulate the dangers of the models that exist today and regulation for the future is doomed to fail. They say that powerful models could be decades away, and it's unreasonable to regulate that far in advance. What do you think? Should regulation proactively prepare for more powerful models?

Yes, regulations should prepare for more powerful models	67%
No, regulations should only focus on models that exist today	13%
Don't know	20%
Totals	100%
N	**1,118**

7. There is a debate on whether regulations should explicitly restrict how much more powerful AI models are allowed to become. Supporters of restricting the power of AI models say that the burden of proof for safety should fall on the AI tech companies. They say that the models being created by AI companies might be excessively risky and dangerous, and simply waiting for dangerous models to be released could be courting disaster. Opponents of restricting the power of AI models say that we can decide on proper regulation when/if models become dangerous. They say that limiting model power is incredibly difficult to do with regulation and doing so risks making American companies fall behind China. What do you think? Should regulation restrict how powerful models can become until safety can be assessed?

Yes, we should restrict how powerful models can become	67%
No, we should not restrict how powerful models can become	14%
Don't know	19%
Totals	100%
N	**1,118**

8. There is a debate on whether regulations should be made with the goal of making it more difficult to release powerful AI.Supporters of such regulations say that more powerful models are currently very risky. They claim that their potential for accidents and misuse range from spreading propaganda to even creating powerful biological weapons. They say stringent restrictions are necessary. Opponents of such regulation say that onerous regulations will only hamper innovation and put the US behind China. If it's made more difficult to release models, consumers won't benefit from the technology. They also claim that the potential for misuse is overblown. What do you think? Should there be regulations to make it more difficult to release powerful AI?

Yes, regulations should have the goal of making it more difficult to release powerful AI	66%
No, regulations should not have the goal of making it more difficult to release powerful AI	14%
Don't know	20%
Totals	100%
N	**1,118**

9. There is a debate over whether regulation of AI should restrict current state of the art models. Supporters of regulating current models say that they are already risky. They claim that even current models can be used to spread propaganda, be used to learn how to make bombs, and more. They say that dealing with current harm effectively will also prevent worse harm with more powerful models. Opponents of regulating current models say that the models bring far greater benefits than their minimal risks. They claim that the models already have extensive safety protocols and further regulation would stifle innovation for little concrete benefit. They say these models are too weak to be restricted. What do you think? Should we regulate current state of the art AI models?

Yes, we should regulate current state of the art AI models	66%
No, we should not regulate current state of the art AI models	14%
Don't know	20%
Totals	100%
N	**1,118**

10. There is currently a debate about how regulations around AI should be structured, whether it should treat AI as a dangerous powerful technology with growing capabilities or as a piece of consumer technology with real but limited impact. Supporters of regulating AI as a dangerous powerful technology claim that the power of AI is growing rapidly. AI can already write at a college level, and soon it could develop dangerous technology, automate massive amounts of human work, and transform society. They say we need to treat it as the powerful technology it is. Opponents of regulating AI as a dangerous powerful technology claim that the threat of AI is overblown. They say it is merely a piece of consumer technology. It will impact society, but similar to how smartphones impact society. They claim that treating it as dangerous is fear mongering, and will lead to excessive counterproductive regulation. What do you think? Should AI be regulated like a dangerous powerful technology?

Yes, AI should be regulated like a dangerous powerful technology	69%
No, AI should not be regulated like a dangerous powerful technology	15%
Don't know	16%
Totals	100%
N	**1,118**

11. There is currently a debate about whether regulation should have the goal of delaying super intelligence. Super intelligent AI is AI that vastly exceeds the capabilities of humans. Supporters of regulation to delay super intelligence say that the AI labs are actively pursuing superhuman AI that is more powerful than humans. They claim that such AI is very dangerous with unknown consequences. It could lead to catastrophic risks to humanity. They say super intelligence hasn't been proven safe and we aren't ready. Opponents of regulation to delay super intelligence say that super intelligence is a long time away, if it will ever arrive. They claim it's fearmongering to worry about this. Additionally, they claim that when superintelligence does arrive, it will be beneficial for humanity, able to advance research to cure diseases. What do you think? Should regulation take the active goal of preventing AI superintelligence?

Yes, regulation should aim to actively prevent AI superintelligence	63%
No, regulation shouldn't aim to actively prevent AI superintelligence	16%
Don't know	21%
Totals	100%
N	**1,118**

12. There is a debate among policy makers on whether it's more important to regulate near term weaker threats from AI or dangerous but unknown threats. Those who favor targeting near term weaker threats say that AI is already causing problems from misinformation to bias. They claim that we can only regulate problems that exist today and trying to deal with unknown problems is a fool's errand. Those who favor targeting more dangerous but unknown threats say that the biggest dangers of AI are ahead of us. AI is advancing rapidly, and the dangers could be catastrophic. We can require stringent safety checks before release for even unknown future models. What do you think? Should we focus regulation on weaker but known threats today, or on unknown but more dangerous threats in the future?

Weaker but known threats	22%
More dangerous but unknown threats	46%
Don't know	32%
Totals	100%
N	**1,118**

13. Lately, there has been debate around the timing of AI regulation. Some policymakers want to move quickly and restrict dangerous models now while other policymakers want to wait and take our time regulating. Supporters of regulating quickly say that AI models are already getting powerful and are about to become much more powerful. AI companies are moving full speed ahead to develop models which could cause unknown and potentially catastrophic risks to humanity while upending our economy. They say we can't afford to wait. Opponents of regulating quickly say that careful regulation takes time. They say that to rush into regulating AI would be misguided. By regulating too fast, we could have the effect of stifling innovation and could allow adversaries like China to get the edge on us. We need to take our time and respond to models as they develop. What do you think? Should we regulate dangerous models quickly or wait and take our time to regulate?

We should regulate potentially dangerous models quickly	62%
We should wait and take our time to regulate	19%
Don't know	19%
Totals	100%
N	**1,118**

14. There is a proposal in Congress that expands access to AI education and hands-on research and training opportunities for all students, including those in community colleges and state universities, thus leveling the educational field and fostering innovation. Do you [support or oppose] such a proposal?

Strongly support	19%
Somewhat support	36%
Somewhat oppose	13%
Strongly oppose	11%
Don't know	21%
Totals	100%
N	**1,118**

15. Should companies heavily relying on AI data analysis be subjected to high-level regulation to prevent privacy violations?

Should be subjected to high-level regulations	66%
Shouldn't be subjected to high-level regulations	14%
Don't know	20%
Totals	100%
N	**1,118**

16. Do you [support or oppose] a law that requires all political advertisements using AI-generated imaging or voice replication to clearly disclose it in advance of the 2024 election?

Strongly support	62%
Somewhat support	16%
Somewhat oppose	6%
Strongly oppose	5%

Don't know	12%
Totals	101%
N	**1,118**

17. Do you [support or oppose] the government funding research that can be used to protect online data and secrets from cyber attack?

Strongly support	50%
Somewhat support	30%
Somewhat oppose	5%
Strongly oppose	4%
Don't know	11%
Totals	100%
N	**1,118**

18. Imagine the election for Congress was taking place in your district today. Which candidate would you vote for?

A candidate who says that artificial intelligence will be a top priority for them. They will focus on policies to ensure artificial intelligence does not threaten American jobs, create a deadly pandemic or accelerate war.	58%
A candidate who says that artificial intelligence will not be a top priority for them. They believe policymakers shouldn't focus on fancy new technologies and should focus on other high-priority issues instead.	42%
Totals	100%
N	**1,118**

19. How often do you use/watch TikTok?

Very often	12%
Somewhat often	13%
Very little	15%
I do not use TikTok	60%
Totals	100%
N	**1,118**

20. When you see a TikTok creator with a large following – in the millions for example – is your assumption that the followers are real, or that some could be AI-generated?

Shown to respondents who use TikTok

They are real	27%

Some could be AI-generated	52%
Don't know	21%
Totals	100%
N	**447**

21. How concerned are you that AI could cause a spike in fake followers/fake users on your favorite social media platform?

Very concerned	28%
Somewhat concerned	27%
Slightly concerned	12%
Not at all concerned	19%
Don't know	14%
Totals	100%
N	**1,118**

22. Some policymakers have suggested that the government should monitor the usage of what are known as "large computer clusters," which are data centers filled with computer chips used to train AI models like ChatGPT.Supporters of this policy say that large compute clusters can be used to create powerful and dangerous AI models, for example to automate hacking or to design new bioweapons. They say that monitoring usage is a key first step to appropriate oversight of the industry and that early monitoring reduces the risk of over-regulation. Opponents of this policy say that monitoring compute clusters would put an excessive burden on companies and that it could harm the competitiveness of US companies. Further, they say monitoring of compute is just the first step to more onerous regulations. What do you think? Should we monitor the usage of large compute clusters?

Yes, the government should monitor the usage of large compute clusters.	62%
No, the government should let companies do what they want with compute clusters.	12%
Don't know	26%
Totals	100%
N	**1,118**

23. Some have argued that Congress should pass a bill restricting Chinese companies' access to US cloud compute services from companies like Google and Amazon. Supporters argue that this is the right next step following export controls on computing hardware that restricted Chinese companies' access to advanced chips implemented by the Biden administration last year. They say Chinese companies can get around the restrictions by purchasing compute through Google or Amazon's web platforms. Detractors say that limiting access to cloud compute would force Chinese companies to rely on local hardware, accelerating the development of Chinese chip manufacturing capabilities. They also argue that China is a key business partner and that restricting access to cloud compute would further damage business relationships between the US and China. What do you think? Should we restrict Chinese companies' access to cloud compute?

Yes, we should restrict Chinese companies' access to cloud compute from US companies.	73%

No, we should let Chinese companies continue to access cloud compute from US companies. 9%
Don't know 18%
Totals 100%
N **1,118**

24. Some policymakers have suggested that the creators of AI models should be liable for harms caused by the models they release. Supporters of this policy say tech companies will take too many risks and release unsafe technologies unless they are held responsible for the damages from the technologies they create. They say companies have a strong financial incentive to downplay and conceal the risks. Opponents say that the party causing the harm should be liable, not the creator of the technology. For example, Dell isn't liable for a hacker using their laptop for illegal hacking. Holding the creators of technology liable would slow down innovation and investment and seriously harm the competitiveness of the US AI industry. What do you think? Should AI companies being held liable for the harms from the technologies they create?

Yes, AI companies should be liable for harms from technologies they create. 73%
No, AI companies should not be liable for harms from technologies they create. 11%
Don't know 16%
Totals 100%
N **1,118**

25. Some policy makers are proposing that the US government should create a regulatory organization to conduct audits on AI models for safety. They say that we must ensure that any model released is safe and that AI companies can't be trusted to regulate their own models. Other policy makers disagree. They say that the government isn't capable of conducting appropriate auditing of AI models. They say that AI companies have the expertise to handle it internally and must be left to self regulate. What do you think? Should the government set up an AI auditing organization?

Yes, the government should set up an AI auditing organization. 64%
No, the government should not set up an AI auditing organization. 14%
Don't know 22%
Totals 100%
N **1,118**

26. When AI companies create new AI models, they create the models through a process known as a "training run".Some policy makers are proposing that AI companies be required to register large training runs with the government in advance of conducting them. They say that it's essential the public knows when large, potentially dangerous, models are created. Other policy makers disagree. They say that such registration would hamper research, making it difficult for smaller AI companies to innovate. They also claim that such a regulatory regime would be expensive and wasteful. What do you think? Should companies be required to register large training runs?

Yes, AI companies should be required to register large training runs. 66%
No, AI companies should not be required to register large training runs. 11%

Don't know	23%
Totals	100%
N	**1,118**

27. Some policy makers proposed requiring AI companies to undergo testing and certification for powerful AI models, with more stringent testing requirements for more powerful models. They say that allowing AI companies to release powerful models without stringent testing is dangerous. The more powerful the model, the more important stringent testing is. Other policy makers disagree. They say that the proposed system would add bureaucratic overhead, making it harder for small companies to compete. They say that the government doesn't understand what testing is needed, and this policy will just make testing less effective. What do you think? Should AI companies be required to test and certify powerful AI models?

Yes, AI companies should be required to test and certify powerful AI models.	74%
No, AI companies should not be required to test and certify powerful AI models.	11%
Don't know	16%
Totals	101%
N	**1,118**

AIPI AI Survey
Wave 2

YouGov

This survey is based on 1,118 interviews conducted by YouGov on the internet of registered voters. The sample was weighted according to gender, age, race/ethnicity, education, and U.S. Census region based on voter registration lists, the U.S. Census American Community Survey, and the U.S. Census Current Population Survey, as well as 2020 Presidential vote. Respondents were selected from YouGov to be representative of registered voters. The weights range from 0.27 to 3.24 with a mean of 1 and a standard deviation of 0.4.

The margin of error (a 95% confidence interval) for a sample percentage p based upon the subsetted sample is approximately 3.2%. It is calculated using the formula:

$$\hat{p} \pm 100 \times \sqrt{\frac{1 + \text{CV}^2}{n}}$$

where CV is the coefficient of variation of the sample weights and n is the sample size used to compute the proportion. This is a measure of sampling error (the average of all estimates obtained using the same sample selection and weighting procedures repeatedly). The sample estimate should differ from its expected value by less than margin of error in 95 percent of all samples. It does not reflect non-sampling errors, including potential selection bias in panel participation or in response to a particular survey.Some items held for future release.